

## Deep Learning Customer Churn

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

```
df=pd.read_excel("Churn_Modelling.xlsx")
```

```
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	Estim
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	0	1
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	1	0
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	1



```
df.info()
```

Saving... Frame ' > 5 9999

#	Column	Non-Null Count	Dtype	
0	RowNumber	10000	non-null	int64

```

1 CustomerId      10000 non-null  int64
2 Surname        10000 non-null  object
3 CreditScore    10000 non-null  int64
4 Geography      10000 non-null  object
5 Gender         10000 non-null  object
6 Age            10000 non-null  int64
7 Tenure         10000 non-null  int64
8 Balance        10000 non-null  float64
9 NumOfProducts  10000 non-null  int64
10 HasCrCard     10000 non-null  int64
11 IsActiveMember 10000 non-null  int64
12 EstimatedSalary 10000 non-null  float64
13 Exited        10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB

```

`df.describe()`

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
<b>count</b>	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
<b>mean</b>	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	10000.000000
<b>std</b>	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499797	50000.000000
<b>min</b>	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000	0.000000
<b>25%</b>	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000	50000.000000
<b>50%</b>	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000000	100000.000000
<b>75%</b>	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.000000	140000.000000
<b>max</b>	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000	1.000000	190000.000000



Saving...



`df.isnull().sum()`

```
RowNumber      0  
CustomerId     0  
Surname        0  
CreditScore    0  
Geography      0  
Gender          0  
Age             0  
Tenure          0  
Balance         0  
NumOfProducts   0  
HasCrCard       0  
IsActiveMember  0  
EstimatedSalary 0  
Exited          0  
dtype: int64
```

```
df.drop(columns=["RowNumber", "CustomerId", "Surname"], inplace=True)
```

```
df.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited	⊕
0	619	France	Female	42	2	0.00		1	1	1	101348.88	1
1	608	Spain	Female	41	1	83807.86		1	0	1	112542.58	0
2	502	France	Female	42	8	159660.80		3	1	0	113931.57	1
3	699	France	Female	39	1	0.00		2	0	0	93826.63	0
4	850	Spain	Female	43	2	125510.82		1	1	1	79084.10	0

### ▼ Target Column

Saving...



### ▼ Exited

```
df['Exited'].describe()
```

```
count    10000.000000
mean      0.203700
std       0.402769
min      0.000000
25%     0.000000
50%     0.000000
75%     0.000000
max      1.000000
Name: Exited, dtype: float64
```

```
df['Exited'].value_counts(normalize=True)
```

```
0    0.7963
1    0.2037
Name: Exited, dtype: float64
```

```
df['Exited'].value_counts().plot(kind='pie', autopct="%1.2f%%", labels=['Retained', 'Exited'], explode=[0, 0.2], shadow=True)
plt.title("Proportion of customer churned and retained", size = 20)
plt.show()
```

Saving...



# Proportion of customer churned and retained



## ▼ Categorical Features

```
#function is created to explore Feature of datatype=Object

def feature_info(feature):
    print(df[feature].value_counts())
    print('\n')
    print(df.groupby(feature)['Exited'].value_counts(normalize=True)*100)

    ax=sns.countplot(data=df,x=feature,hue='Exited')
    plt.title(f"{feature} wise numbers of customer churned and retained",size = 20)
    plt.legend(labels=['0 : Reatained','1 : Churned'])
    for container in ax.containers:
        ax.bar_label(container)

    return feature
```

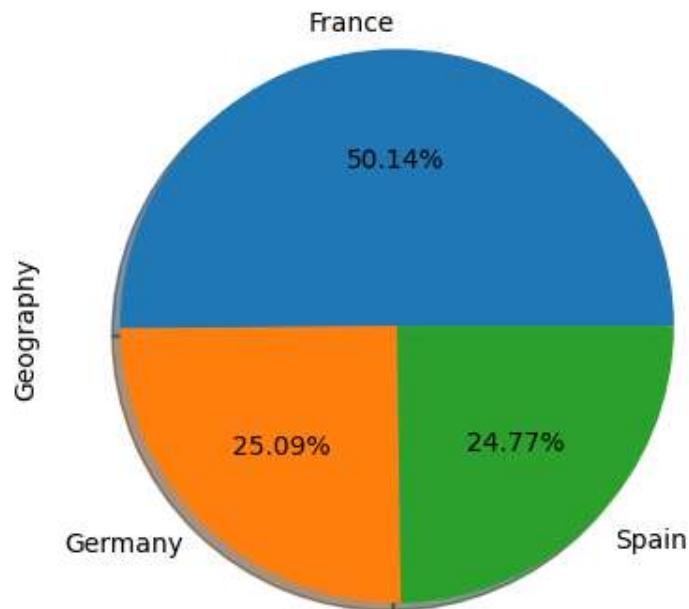
## ▼ Geography

```
df['Geography'].value_counts()
```

France	5014
Saving...	X

```
df[ 'Geography' ].value_counts().plot(kind='pie', autopct="%1.2f%%", shadow=True)
plt.title("Portion of customer geography-wise", size = 20)
plt.show()
```

## Portion of customer geography-wise



```
feature_info("Geography")
```

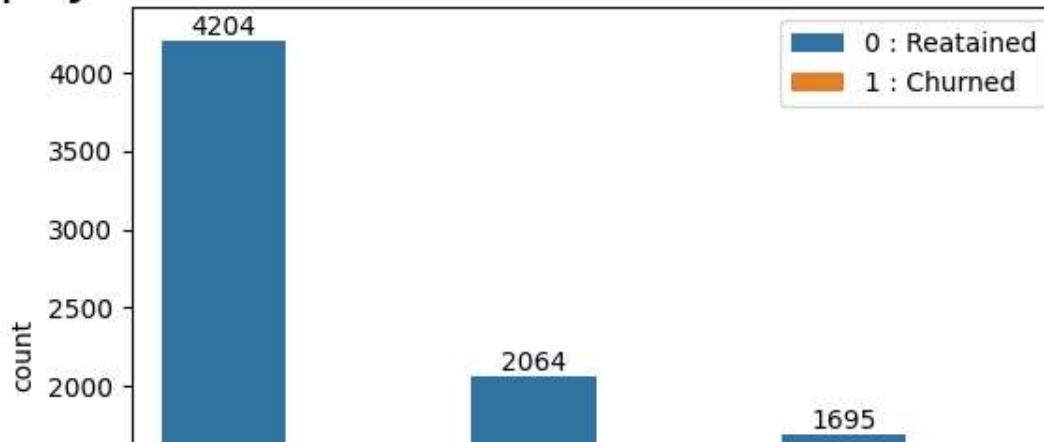
Saving...

X

```
France      5014
Germany    2509
Spain       2477
Name: Geography, dtype: int64
```

```
Geography  Exited
France      0        83.845233
              1        16.154767
Germany    0        67.556796
              1        32.443204
Spain       0        83.326605
              1        16.673395
Name: Exited, dtype: float64
'Geography'
```

## Geography wise numbers of customer churned and retained



## ▼ Gender

```
df['Gender'].describe()
```

Saving...

X

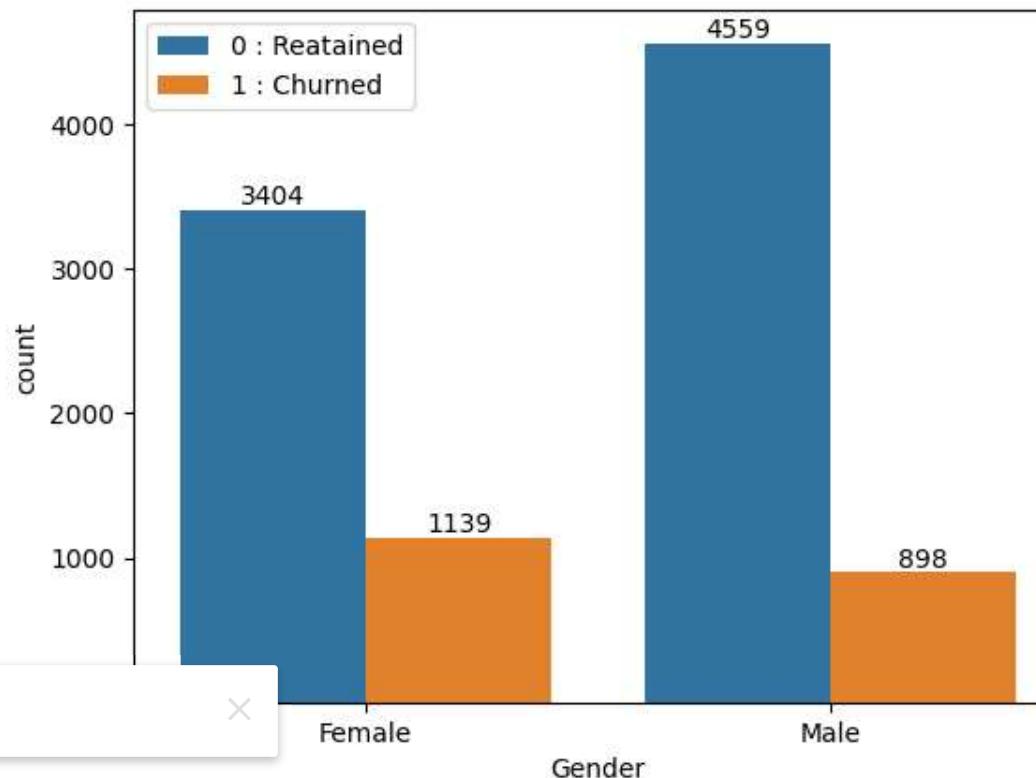
```
top      Male
freq     5457
Name: Gender, dtype: object
```

```
feature_info('Gender')
```

```
Male      5457  
Female    4543  
Name: Gender, dtype: int64
```

```
Gender  Exited  
Female  0           74.928461  
        1           25.071539  
Male    0           83.544072  
        1           16.455928  
Name: Exited, dtype: float64  
'Gender'
```

## Gender wise numbers of customer churned and retained



## ▼ Number of Product

```
feature_info('NumOfProducts')
```

Saving...



```
1    5084
2    4590
3     266
4      60
Name: NumOfProducts, dtype: int64
```

NumOfProducts\_Fxited

## ▼ Has Credit Card or Not

0: No Credit Card

1: Has Credit Card

Name: Fxited dtype: float64

```
df['HasCrCard'].value_counts().plot(kind='pie', autopct="%1.2f%%", labels=['Has Credit Card', 'No Credit Card'], shadow=True)
plt.title("Proportion of customer have credit or Not", size = 20)
plt.legend()
plt.show()
```

Saving...



# Proportion of customer have credit or Not

```
feature_info('HasCrCard')
```

Saving...



```
1      7055  
^      2215
```

## ▼ Tenure

```
HasCrCardExited  
feature_info('Tenure')
```

Saving...



```
    ..  
10    490  
0     413  
Name: Tenure, dtype: int64
```

```
Tenure  Exited  
0       0      76.997579  
       1      23.002421  
1       0      77.584541  
       1      22.415459  
2       0      80.820611  
       1      19.179389  
3       0      78.889990  
       1      21.110010  
4       0      79.474216  
       1      20.525784  
5       0      79.347826  
       1      20.652174  
6       0      79.731127  
       1      20.268873  
7       0      82.782101  
       1      17.217899  
8       0      80.780488  
       1      19.219512  
9       0      78.353659  
       1      21.646341  
10      0      79.387755  
        1      20.612245  
Name: Exited, dtype: float64
```

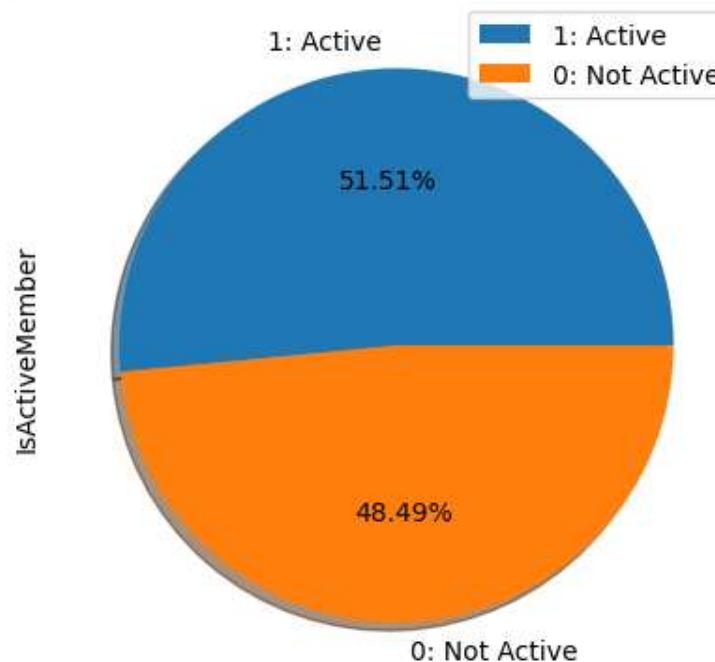
'Tenure'

Tenure wise numbers of customer churned and retained

## ▼ Customer Active Or Not

```
df['IsActiveMember'].value_counts().plot(kind='pie', autopct="%1.2f%%", labels=['1: Active','0: Not Active'], shadow=True)  
                                         title="Customer Active or Not", size = 20)  
  
Saving...  
pic.show()
```

## Proportion of customer Active or Not



```
feature_info('IsActiveMember')
```

Saving...

X

```
1    5151
0    4849
Name: IsActiveMember, dtype: int64
```

```
IsActiveMember  Exited
0              0    73.149103
                1    26.850897
1              0    85.730926
                1    14.269074
```

```
Name: Exited, dtype: float64
```

```
'IsActiveMember'
```

## IsActiveMember wise numbers of customer churned and retained



## ▼ Numerical Feature

```
def Num_feature(feature):
    print(df[feature].describe())
    fig, axes = plt.subplots(2, 2, figsize=(15, 10))
    sns.scatterplot(x=df[feature], y=df['Exited'], hue=df['Exited'], ax=axes[0][0])
    axes[0][0].set_title((f'Scatter plot {feature} VS Customer Churn'))
    sns.boxplot(y=df[feature], x=df['Exited'], ax=axes[0][1])
    axes[0][1].set_title((f'Distribution of {feature} VS Customer Churn'))
    #axes[0][1].legend(labels=['0 : Retained', '1 : Churned'], loc='center')

    sns.distplot(x=df[feature], kde=True, ax=axes[1][0])
    axes[1][0].set_title((f'Distribution of {feature}'))
```

Saving...

```
sns.boxplot(y=df[feature], ax=axes[1][1])
    axes[1][1].set_title((f'Distribution of {feature} using boxPlot'))
    fig.subplots_adjust(wspace=0.2, hspace=0.3)
```

## ▼ Age

```
Num_feature('Age')
```

Saving...



```
count    10000.000000
mean     38.921800
std      10.487806
min      18.000000
25%     32.000000
50%     37.000000
75%     44.000000
max      92.000000
Name: Age, dtype: float64
```

Scatter plot Age VS Customer Churn



Distribution of Age VS Customer Churn



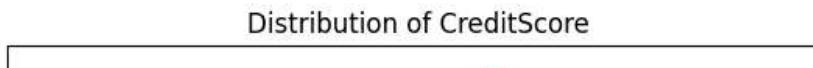
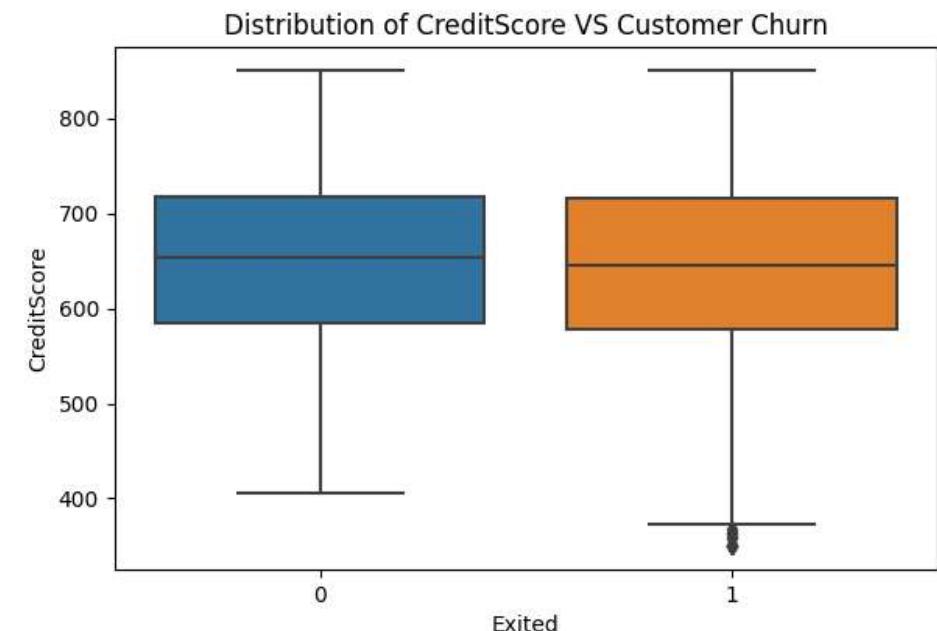
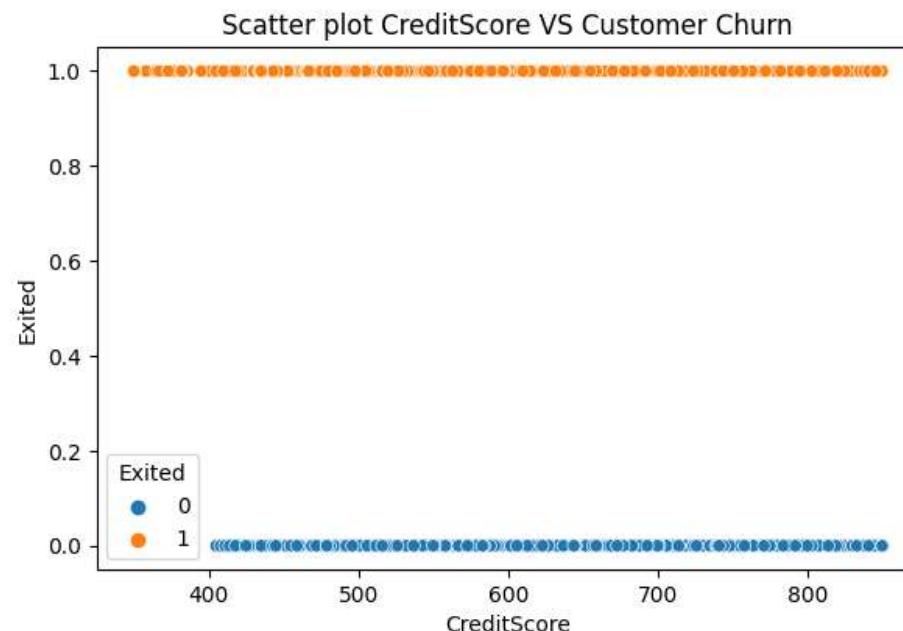
## ▼ Credit Score

```
Num_feature('CreditScore')
```

Saving...



```
count    10000.000000
mean     650.528800
std      96.653299
min     350.000000
25%    584.000000
50%    652.000000
75%    718.000000
max     850.000000
Name: CreditScore, dtype: float64
```



## ▼ Balance

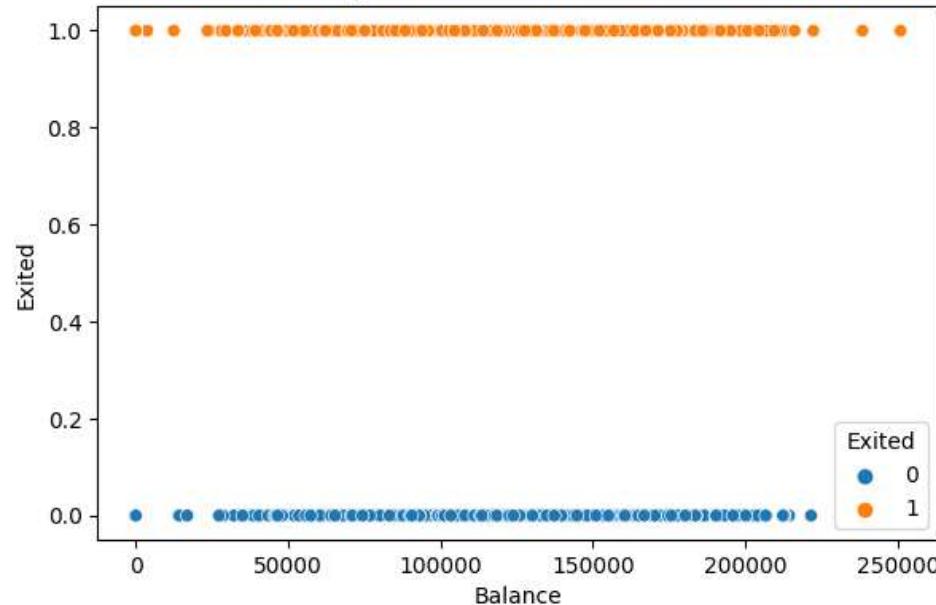
Saving... X

```

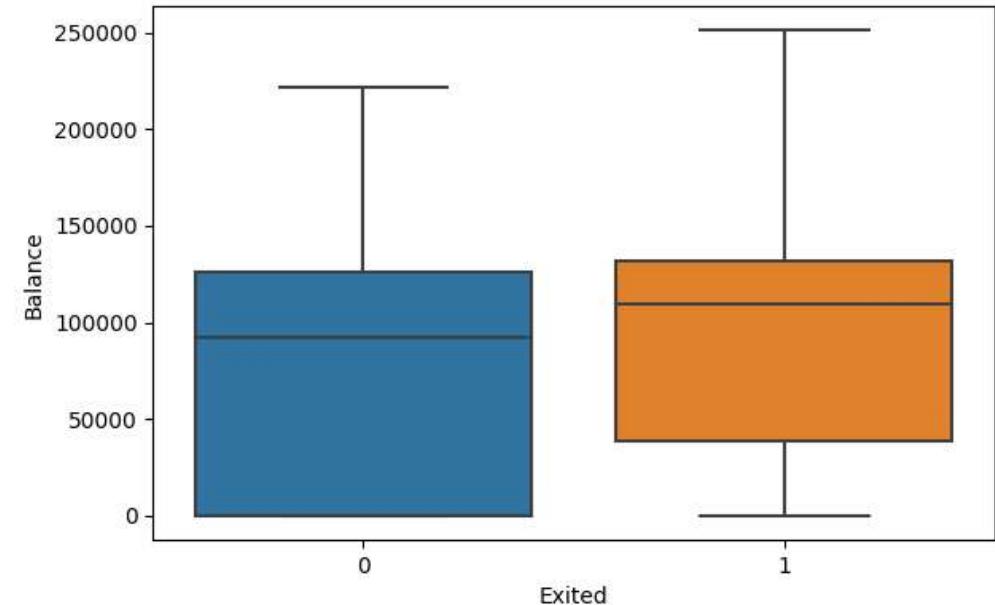
count    10000.000000
mean     76485.889288
std      62397.405202
min      0.000000
25%     0.000000
50%     97198.540000
75%    127644.240000
max    250898.090000
Name: Balance, dtype: float64

```

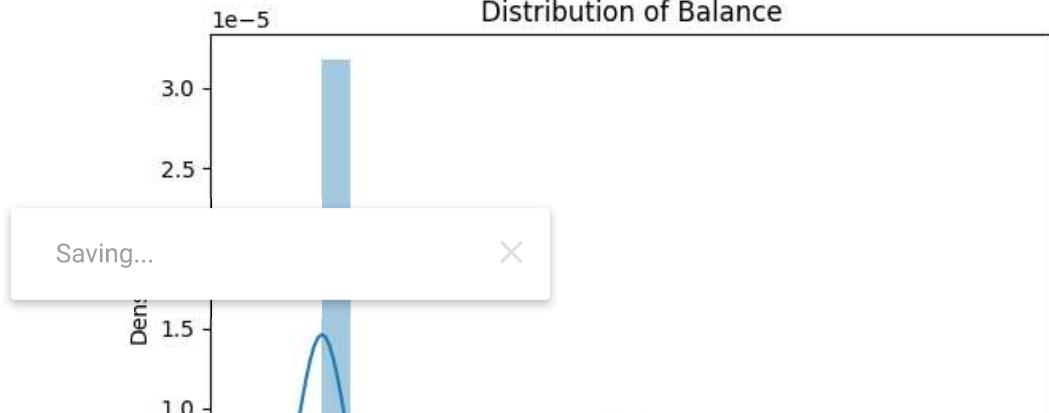
Scatter plot Balance VS Customer Churn



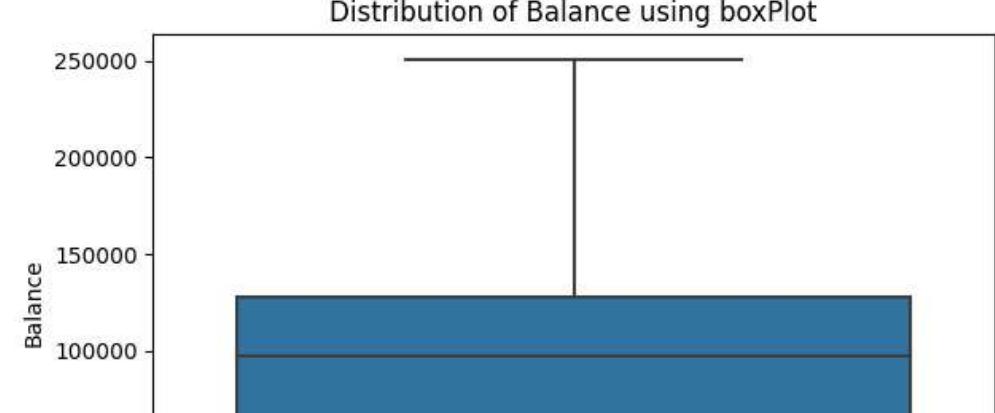
Distribution of Balance VS Customer Churn

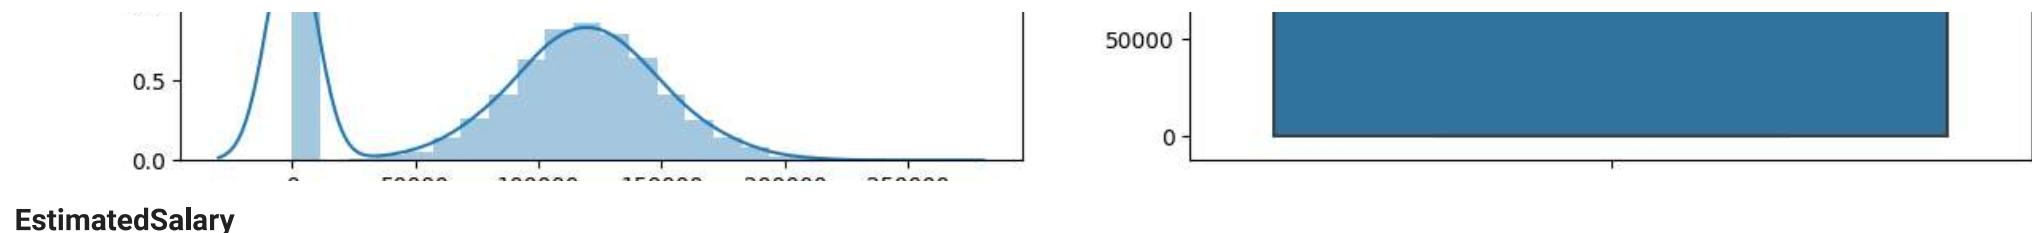


Distribution of Balance



Distribution of Balance using boxPlot



**EstimatedSalary**

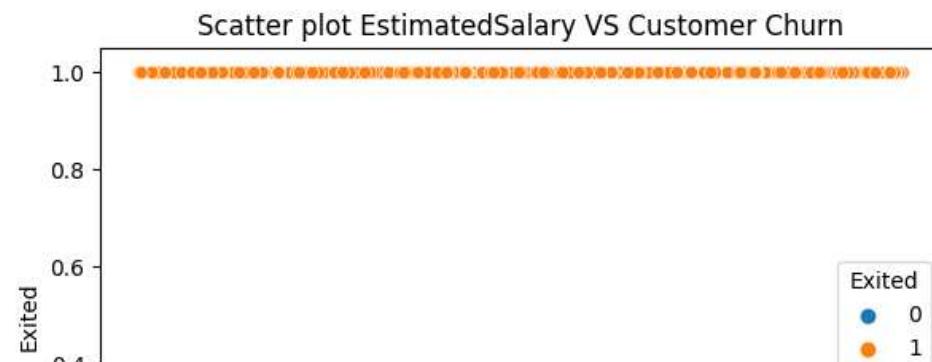
```
Num_feature('EstimatedSalary')
```

Saving...

X

```
count    10000.000000
mean     100090.239881
std      57510.492818
min      11.580000
25%     51002.110000
50%     100193.915000
75%     149388.247500
max     199992.480000
```

Name: EstimatedSalary, dtype: float64

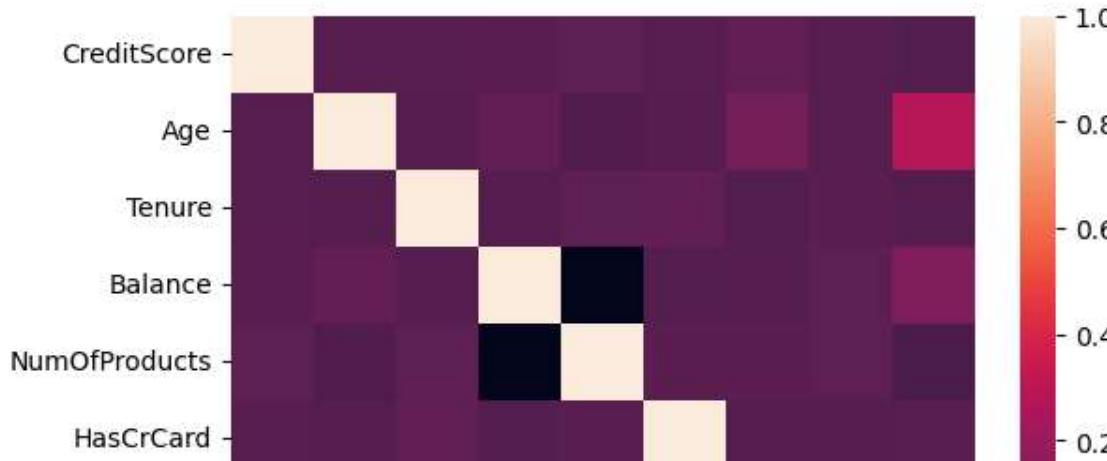


## ▼ Correlation

```
sns.heatmap(df.corr())
```

Saving...

&lt;Axes: &gt;



```
df.corr()['Exited'].sort_values()
```

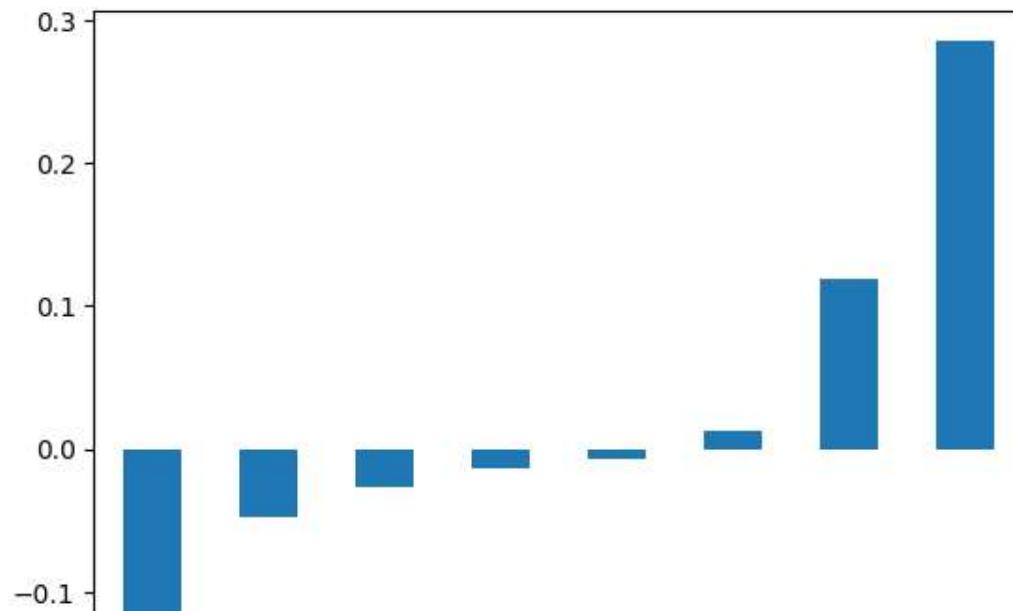
IsActiveMember	-0.156128
NumOfProducts	-0.047820
CreditScore	-0.027094
Tenure	-0.014001
HasCrCard	-0.007138
EstimatedSalary	0.012097
Balance	0.118533
Age	0.285323
Exited	1.000000

Name: Exited, dtype: float64

```
df.corr()['Exited'][:-1].sort_values().plot(kind='bar')
```

Saving...

&lt;Axes: &gt;



## ▼ Encoding

Ex Tr O E ar ar IC Ag

```
from sklearn.preprocessing import LabelEncoder  
le=LabelEncoder()  
df["Geography"]=le.fit_transform(df["Geography"])  
le1=LabelEncoder()  
df["Gender"]=le1.fit_transform(df["Gender"])
```

```
x=df.drop('Exited',axis=1)  
y=df['Exited'].values
```

Saving...

X

```
from sklearn.preprocessing import StandardScaler  
  
std=StandardScaler()  
  
x=std.fit_transform(x)  
x  
  
array([[-0.32622142, -0.90188624, -1.09598752, ..., 0.64609167,  
       0.97024255, 0.02188649],  
      [-0.44003595, 1.51506738, -1.09598752, ..., -1.54776799,  
       0.97024255, 0.21653375],  
      [-1.53679418, -0.90188624, -1.09598752, ..., 0.64609167,  
       -1.03067011, 0.2406869 ],  
      ...,  
      [ 0.60498839, -0.90188624, -1.09598752, ..., -1.54776799,  
       0.97024255, -1.00864308],  
      [ 1.25683526, 0.30659057, 0.91241915, ..., 0.64609167,  
       -1.03067011, -0.12523071],  
      [ 1.46377078, -0.90188624, -1.09598752, ..., 0.64609167,  
       -1.03067011, -1.07636976]])
```

## ▼ Train Test Split

```
from sklearn.model_selection import train_test_split  
xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.30,random_state=42,stratify=y)
```

```
import tensorflow as tf  
from tensorflow.keras import Sequential  
Saving...  
    × Dense  
        classification_report,confusion_matrix
```

```
ann=Sequential()
ann.add(Dense(units=20,activation="relu"))
ann.add(Dense(units=20,activation="relu"))
ann.add(Dense(units=1,activation="sigmoid"))

ann.compile(optimizer="adam",loss="binary_crossentropy")

ann.fit(xtrain,ytrain,epochs=200,batch_size=32,validation_data=(xtest, ytest))
```

```
Epoch 1/200
219/219 [=====] - 7s 4ms/step - loss: 0.6295 - val_loss: 0.4669
Epoch 2/200
219/219 [=====] - 1s 4ms/step - loss: 0.4438 - val_loss: 0.4113
Epoch 3/200
219/219 [=====] - 1s 5ms/step - loss: 0.4142 - val_loss: 0.3884
Epoch 4/200
219/219 [=====] - 1s 6ms/step - loss: 0.3942 - val_loss: 0.3710
Epoch 5/200
219/219 [=====] - 1s 4ms/step - loss: 0.3800 - val_loss: 0.3644
Epoch 6/200
219/219 [=====] - 1s 4ms/step - loss: 0.3703 - val_loss: 0.3572
Epoch 7/200
219/219 [=====] - 1s 4ms/step - loss: 0.3637 - val_loss: 0.3517
Epoch 8/200
219/219 [=====] - 1s 4ms/step - loss: 0.3587 - val_loss: 0.3510
Epoch 9/200
219/219 [=====] - 1s 4ms/step - loss: 0.3558 - val_loss: 0.3457
Epoch 10/200
219/219 [=====] - 1s 4ms/step - loss: 0.3532 - val_loss: 0.3429
Epoch 11/200
219/219 [=====] - 1s 4ms/step - loss: 0.3507 - val_loss: 0.3418
Epoch 12/200
219/219 [=====] - 1s 4ms/step - loss: 0.3486 - val_loss: 0.3407
Epoch 13/200
219/219 [=====] - 1s 4ms/step - loss: 0.3472 - val_loss: 0.3401
Epoch 14/200
219/219 [=====] - 1s 4ms/step - loss: 0.3454 - val_loss: 0.3406
Epoch 15/200
219/219 [=====] - 1s 4ms/step - loss: 0.3444 - val_loss: 0.3377
Epoch 16/200
219/219 [=====] - 1s 4ms/step - loss: 0.3425 - val_loss: 0.3374
```

```
Epoch 17/200
219/219 [=====] - 1s 6ms/step - loss: 0.3417 - val_loss: 0.3370
Epoch 18/200
219/219 [=====] - 1s 5ms/step - loss: 0.3405 - val_loss: 0.3340
Epoch 19/200
219/219 [=====] - 1s 4ms/step - loss: 0.3395 - val_loss: 0.3382
Epoch 20/200
219/219 [=====] - 1s 4ms/step - loss: 0.3388 - val_loss: 0.3381
Epoch 21/200
219/219 [=====] - 1s 4ms/step - loss: 0.3371 - val_loss: 0.3359
Epoch 22/200
219/219 [=====] - 1s 4ms/step - loss: 0.3368 - val_loss: 0.3345
Epoch 23/200
219/219 [=====] - 1s 4ms/step - loss: 0.3361 - val_loss: 0.3381
Epoch 24/200
219/219 [=====] - 1s 5ms/step - loss: 0.3345 - val_loss: 0.3340
Epoch 25/200
219/219 [=====] - 1s 4ms/step - loss: 0.3339 - val_loss: 0.3328
Epoch 26/200
219/219 [=====] - 1s 4ms/step - loss: 0.3329 - val_loss: 0.3359
Epoch 27/200
219/219 [=====] - 1s 4ms/step - loss: 0.3322 - val_loss: 0.3342
Epoch 28/200
219/219 [=====] - 1s 4ms/step - loss: 0.3312 - val_loss: 0.3323
Epoch 29/200
.../.../.../...
```

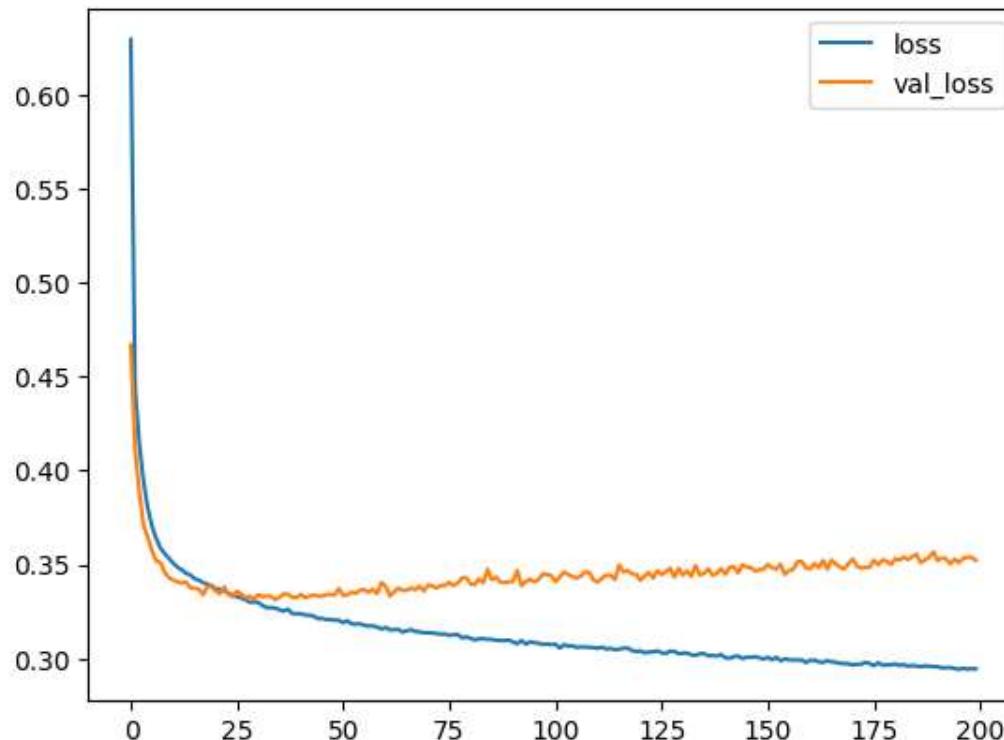
```
#ann.history.history
```

```
model_loss = pd.DataFrame(ann.history.history)
model_loss.head()
```

	loss	val_loss	edit
0	0.629473	0.466923	
1	0.443821	0.411328	
2	0.394162	0.371012	
3	0.379991	0.364356	

```
model_loss.plot()
```

<Axes: >



```
ypred=ann.predict(xtest)
```

94/94 [=====] - 0s 2ms/step

```
ypred
```

Saving...

X

```
[0.05692097],  
[0.02869534],  
...,
```

```
[0.5670314 ],  
[0.31318736],  
[0.00165451]], dtype=float32)
```

```
ypred=np.where(ypred>0.5,1,0)  
ypred
```

```
array([[0],  
       [0],  
       [0],  
       ...,  
       [1],  
       [0],  
       [0]])
```

```
print(classification_report(ytest,ypred))
```

	precision	recall	f1-score	support
0	0.89	0.94	0.91	2389
1	0.71	0.53	0.60	611
accuracy			0.86	3000
macro avg	0.80	0.74	0.76	3000
weighted avg	0.85	0.86	0.85	3000

```
print(confusion_matrix(ytest,ypred))
```

```
[[2257 132]  
 [ 289 322]]
```

Saving...



```
ann=Sequential()
ann.add(Dense(units=20,activation="relu"))
ann.add(Dense(units=20,activation="relu"))
ann.add(Dense(units=1,activation="sigmoid"))
ann.compile(optimizer="adam",loss="binary_crossentropy",metrics=['accuracy'])

from tensorflow.keras.callbacks import EarlyStopping

help(EarlyStopping)
```

Saving...

X

```
on_epoch_end(self, batch, logs=None,
    Called at the end of a training batch in `fit` methods.

    Subclasses should override for any actions to run.

    Note that if the `steps_per_execution` argument to `compile` in
    `tf.keras.Model` is set to `N`, this method will only be called every
    `N` batches.

    Args:
        batch: Integer, index of batch within the current epoch.
        logs: Dict. Aggregated metric results up until this batch.

    set_model(self, model)

    set_params(self, params)

-----
Data descriptors inherited from Callback:

__dict__
    dictionary for instance variables (if defined)

__weakref__
    list of weak references to the object (if defined)
```

```
early_stop = EarlyStopping(monitor='val_loss', mode='min', patience=25)
```

```
ann.fit(xtrain,ytrain,epochs=200,batch_size=32,validation_data=(xtest, ytest),callbacks=[early_stop])
```

Saving... X

Saving...



```
Epoch 50/200
219/219 [=====] - 1s 6ms/step - loss: 0.3229 - accuracy: 0.8637 - val_loss: 0.3337 - val_accuracy: 0.8643
<keras.callbacks.History at 0x7fbab83b5240>
```

```
#help(ann.fit)
```

```
model_loss = pd.DataFrame(ann.history.history)
model_loss.head()
```

	loss	accuracy	val_loss	val_accuracy	edit
0	0.523735	0.784714	0.447092	0.810333	
1	0.449816	0.810857	0.414659	0.823000	
2	0.429820	0.819571	0.401808	0.828333	
3	0.414360	0.827714	0.387696	0.834333	
4	0.395488	0.838000	0.369148	0.848333	

```
model_loss[['accuracy','val_accuracy']].plot()
```

Saving...

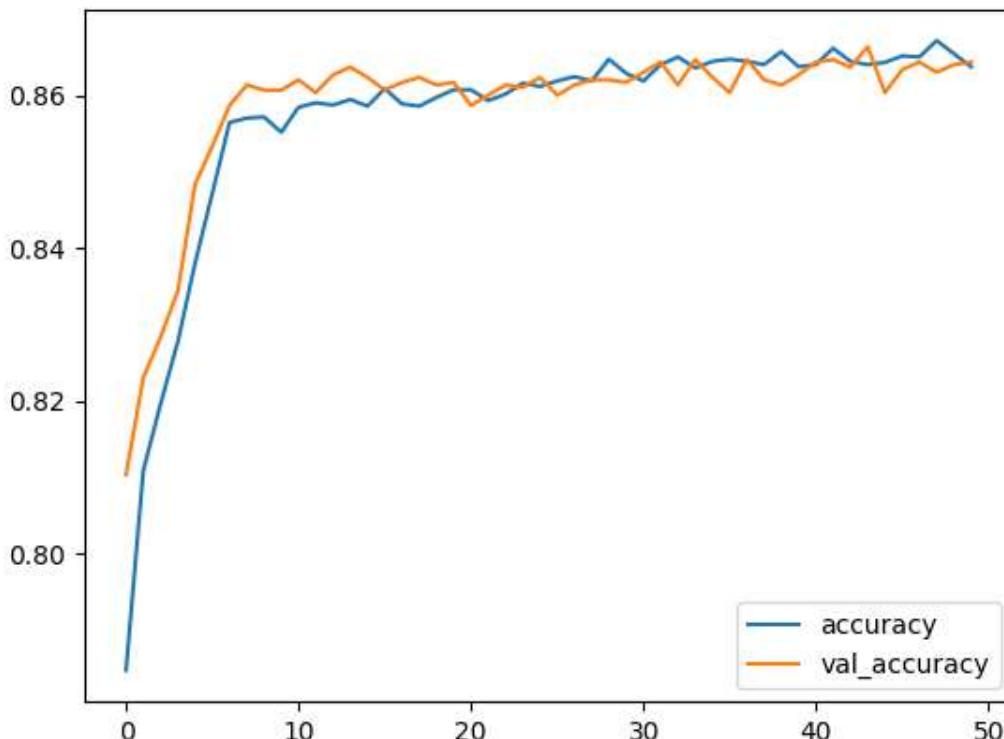


&lt;Axes: &gt;



```
model_loss[['accuracy','val_accuracy']].plot()
```

&lt;Axes: &gt;



```
predictions = ann.predict(xtest)
```

Saving...

=====] - 0s 1ms/step

```
predictions
```

```
array([[0.004102  ],
       [0.11235598],
       [0.05570493],
       ...,
       [0.46079794],
       [0.47469172],
       [0.01805007]], dtype=float32)
```

## ▼ Adding in DropOut Layers

```
from tensorflow.keras.layers import Dropout

ann=Sequential()
ann.add(Dense(units=20,activation="relu"))
ann.add((Dropout(0.5)))

ann.add(Dense(units=20,activation="relu"))
ann.add((Dropout(0.5)))

ann.add(Dense(units=1,activation="sigmoid"))
ann.compile(optimizer="adam",loss="binary_crossentropy",metrics=['accuracy'])

ann.fit(xtrain,ytrain,epochs=200,batch_size=32,validation_data=(xtest, ytest),callbacks=[early_stop])
```

Saving...



Saving...



&lt;keras.callbacks.History at 0x7fa8299960&gt;

```
model_loss = pd.DataFrame(ann.history.history)
model_loss.head()
```

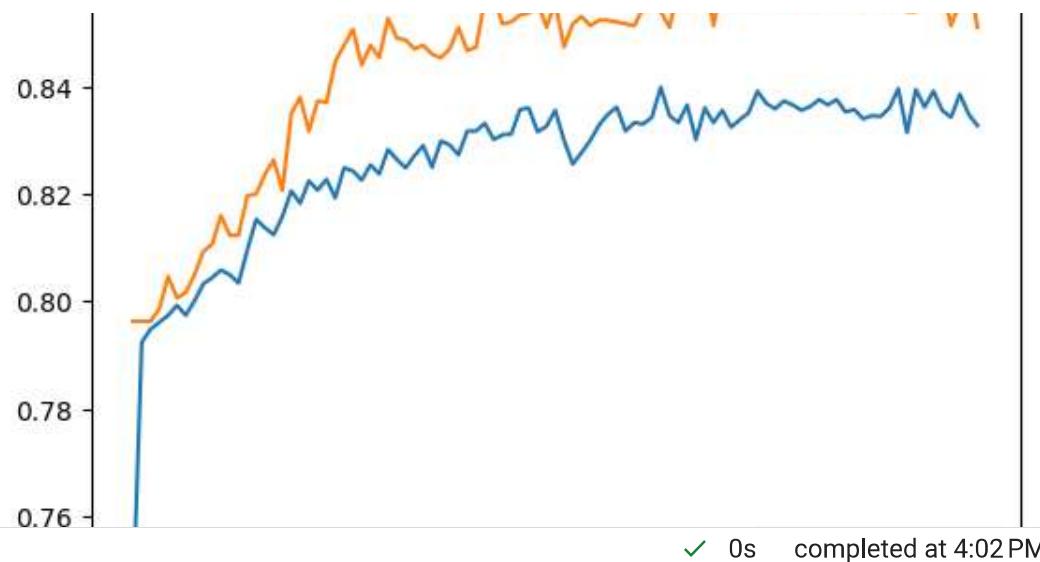
	loss	accuracy	val_loss	val_accuracy	edit
0	0.585261	0.743857	0.490537	0.796333	
1	0.527671	0.792429	0.457221	0.796333	
2	0.497498	0.794857	0.439290	0.796333	
3	0.489286	0.796143	0.431156	0.798667	
4	0.471733	0.797429	0.423902	0.804667	

```
model_loss[['accuracy','val_accuracy']].plot()
```

Saving...



&lt;Axes: &gt;

Saving... X